

Classification and Regression trees

(Recursive partitioning)

R. Bellio & N. Torelli

Spring 2018

University of Udine & University of Trieste

Regression Trees

Classification Trees

Ensemble methods

Regression Trees

Non-parametric regression models

- **Non-parametric** (or semi-parametric) regression modelling keeps the usual specification:

$$y = g(x_1, \dots, x_p, \epsilon)$$

but relaxes the assumption of linearity, and replaces it with a much weaker assumption of a smooth g

- Pro's and con's
 - \mapsto greater flexibility and potentially more accurate estimate of g
 - \mapsto greater computation and often more difficult-to-interpret results: typically used for prediction, not interpretation
- Some examples of nonparametric regression models are:
 - \mapsto Local Polynomial Regression
 - \mapsto Kernel regression
 - \mapsto Smoothing splines
 - \mapsto (Generalized) Additive models
 - \mapsto Decision (regression) trees

Step functions as approximators


- A simple, yet effective, way to approximate a generic function $f(x)$ is to use a step function, that is, a piecewise constant function
- In such a case, there are various choices to be made:
 - where are the subdivision points to be placed?
 - which value of y must be assigned to each interval?
 - how many subdivisions of the x axis must be considered?
- The idea is to generalize the use of step functions to approximate (or predict) a response Y as function of some covariates.
- Note that Y could be of different nature: numeric, factor, count, ...

Step functions as a spline

- A step function actually is a spline of degree 0. Assume we want to fit such a function to a simple set of data.
- Subdivision points are now the knots and their position should be chosen to reflect changes of the function $f(x)$ (for instance more knots where the function is steeper)
- In a given interval the value of the constant can be chosen to be a average of the level of the function itself
- The choice of the number of subdivisions is critical: any increase in the number of steps increases the quality of the approximation, and therefore we are led to think of infinite subdivisions.
- However, this is counter to the requirement to use a approximate representation using few parameters and therefore to adopt a finite number of subdivisions.


An introductory examples

- If y is quantitative a global approximation of y could be its mean. Or we can use a (regression) function $g(\cdot)$




An introductory examples

- Now consider a subdivision on X and approximate y with its local mean \hat{y}_i in the i -th interval and g is a piecewise constant function




The tree

Note that the value \hat{y}_i of the function g can be also described by the following tree




An introductory examples

- As the number of intervals increase, we could achieve a very accurate description of the data




An introductory examples

- As the number of intervals increase, we could achieve a very accurate description of the data (leading to overfitting)



A simple example of tree partitioning for two covariates



In the top right panel first split at $X_1 = t_1$. Then the region $X_1 \leq t_1$ is split at $X_2 = t_2$ and the region $X_1 > t_1$ is split at $X_1 = t_3$. Finally, the region $X_1 > t_3$ is split at $X_2 = t_4$. The result of such a recursive binary splitting is a partition into the five regions R_1, R_2, \dots, R_5 shown in the figure.

- The corresponding regression model predicts Y with a constant c_m in region R_m , that is, $\hat{f}(X_1, X_2) = \sum_{m=1}^5 c_m I\{(X_1, X_2) \in R_m\}$
- The sets R_m are rectangles, in the 2-dimensional space, with their edges parallel to the coordinate axes) and c_1, \dots, c_5 are constants. The top left panel represents a partition that cannot be obtained by recursive binary splitting

A regression tree

- More generally:
 - we want estimate a regression curve $f(x_1, x_2, \dots, x_p)$ underlying the data by $\hat{f}(x_1, x_2, \dots, x_p) = \sum_{m=1}^M c_m I\{(x_1, x_2, \dots, x_p) \in R_m\}$ where $I(x_1, x_2, \dots, x_p \in R_m)$ is the indicator function of the set R_m (R_m are rectangles, in the p -dimensional sense, with their edges parallel to the coordinate axes) and c_1, \dots, c_M are constants.
 - Given an objective function such as the Deviance

$$D = \sum_{i=1}^n (y_i - \hat{f}(x_{1i}, x_{2i}, \dots, x_{pi}))^2$$

- the goal is to define a partition of the space of the covariates that minimizes D

Building the Regression tree

- this minimization, even if we fix the number of the elements of the partition, involves very complex computation
- operatively a sub-optimal approach is considered, of step-by-step optimization: we construct a sequence of gradually more refined approximations and to each of these we minimize the deviance relative to the passage from the current approximation to the previous one.
- It is not ensured that we get the global maximum. This procedure is called greedy-algorithm
- This operation is represented by a series of binary splits
- Each internal node represents a value query on one of the variables – e.g. “Is $x_3 > 0.4$?”. If the answer is ‘Yes’, go right, else go left.
- The terminal nodes are the decision nodes. Typically each terminal node is assigned a value, c_h , given by the arithmetic mean of the observed y_i having component x_{ji} falling in this node.

Growing the tree

- Trees are grown using a random subset of the available data (*the training data*), by recursive splitting
- A terminal node g is split into the left and right daughters (g_L and g_R) that increase the split criterion

$$D_g - D_{g_L} - D_{g_R}$$

the most, where D is the deviance associated to a given node.

- To avoid the overfitting, a large tree T_0 is grown and then pruned backward
- Indeed a tree with n leaves is equivalent to a polynomial regression of degree $n - 1$
- detection of the variable X_j that achieve the best split at each node and which is the split point can be done very quickly and hence by scanning through all of the inputs
- Deviance can be adapted for dealing with a response that is a count or a duration

Pruning the tree

- Pruning criterion: cost of a subtree $T \in T_0$, is defined by

$$C_\alpha(J) = \sum_{j=1}^J D_j + \alpha_j$$

- Here the sum is over the terminal nodes of T , J is the number of terminal nodes in T and α is a cost-complexity parameter
- The choice of an optimal size is evaluated by cross-validation, or on a validation set.
- For each α the best subtree T_α is found via weakest link pruning
- Larger α gives smaller trees
- A best value $\hat{\alpha}$ is estimated via cross-validation (or on a validation set)
- Final chosen tree is $T_{\hat{\alpha}}$
- New observations are classified by passing their x down to a terminal node of the tree, and then using the relative c_h .

An example


The variable FACE refer to the amount of life insurance bought by the head of a household. We want to predict it by using “INCOME”, number of household members, AGE, Education, etc. For illustration, a tree with maximum depth=2 is considered. Package rpart is used.

```
TL <- read.csv("TL.csv", header=TRUE, sep="," , row.names=1)
library(rpart)
attach(TL)
```

```
m2 <- rpart(FACE~INCOME+MARSTAT+NUMHH+EDUCATION+AGE,
            control=rpart.control(maxdepth=2))
m2
```

```
## n= 275
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 275 7.681561e+14  747581.5
##    2) INCOME< 187500 227 2.629158e+14  413511.5
##      4) EDUCATION< 15.5 128 1.075360e+14  239930.5 *
##      5) EDUCATION>=15.5 99 1.465367e+14  637939.4 *
##    3) INCOME>=187500 48 3.600986e+14  2327454.0
##      6) INCOME< 762500 37 1.905974e+14  1870751.0 *
##      7) INCOME>=762500 11 1.358255e+14  3863636.0 *
```

The tree



Regression trees: Advantages

- Logical simplicity and ease of 'communication' (particularly those with a non-quantitative background)
- The step function has a simple, compact mathematical formulation in terms of information to be stored.
- Speed of computation and can take advantage of parallel calculation.
- Can handle huge datasets
- Can handle mixed predictors: quantitative and factors
- Easy ignore redundant variables and automatically detects interactions among variables
- Handle missing data elegantly
- Small trees are easy to interpret

Regression trees: Disadvantages

- Instability of results: very sensitive to the insertion/changes in the sample
- Difficulty in upgrading: if more data arrive, they cannot be added to the already constructed tree; it is necessary to start again from the beginning.
- Difficulty of approximating some mathematically simple functions, particularly if they are steep,
- Statistical inference: formal procedures of statistical inference such as hypothesis testing, confidence intervals, and others are not available.
- (over?) emphasizes interactions
- large trees are hard to interpret
- prediction surface is not smooth

Dealing with missing data

- It is quite common to have observations with missing values for one or more input features. The usual approach is to impute (fill-in) the missing values in some way.
- However, the first issue in dealing with missing data is whether the missing data introduce a sample selection that can bias results.
- It is important to make the distinction between cases where data are
 - Missing Completely at Random (MCAR) mechanism (*no bias*)
 - Missing at Random (MAR) mechanism (*possible bias if the dependence on missingness on some observed covariates are not recognized*)
 - Missing Not at Random (MNAR) (*huge problems, likely to have non negligible selection bias*)
- For the first, and possibly, the second case, in regression trees two approaches can be used when predictors have missing values:
 - if it is categorical, add a specific category for missing values
 - if it is continuous, use surrogate predictors to be used when observation is missing on the primary predictor.

Classification Trees

Classification Trees

- If the target (response) variable is a categorical variable taking values $1, 2, \dots, K$, the only changes needed in the tree algorithm pertain to the criteria for splitting nodes and possibly pruning the tree.
- In these cases the tree will be used for predicting the categorical response and this is labeled as a **classification problem**. And the tree is then a **Classification tree**.
- Also in this case a tree is a hierarchical structure formed by:
 - root: the predictor space
 - nodes:
 1. internal: test an explanatory variable (and splits the predictor space)
 2. terminal (leaf): assign a label class
 - branches: corresponds to values of the explanatory variables
- A tree is constructed by repeated splits of the predictor space (root) into subregions (nodes). Each terminal region is associated with a prediction and their union form a partition of the predictor space.

Growing a classification tree

The following elements are needed

- A set of splits
- A goodness of split criterion
- A stop-splitting rule
- A rule for assigning every terminal node to a class
- Each split depends on the value of a single predictor x_j and depends on the nature of x_j :
 - qualitative, with values in $\mathcal{L} = \{l_1, \dots, l_K\}$: a split is any question as “is $x_j \in S_{\mathcal{L}}$?” with $S_{\mathcal{L}}$ a subset of \mathcal{L} ;
 - quantitative, with range (a, b) : a split is any question as “is $x_j \leq s$?” with $a \leq s < b$
- Examples
 - “Is the age of the subject not greater than 60?”
 - “Is the weather cloudy or rainy?”
- At each step of the tree growing procedure, the best split is identified for each predictor and, among these, the best of the best is selected.

The goodness of split criterion

- The objective of classification tree construction is to finally obtain nodes that are as **pure** as possible, i.e., the split should send towards each branch observations of the same class
- It makes sense to consider good a split when it leads to a high **reduction of impurity** of the node (a high increase of the prediction/classification accuracy).
- Consider a node t for a two class classification problem, the two classes of y have frequency $p(t)$ and $1 - p(t)$. A natural **impurity measure** of a node t is, the so called Misclassification error:

$$i(t) = 1 - \max(p(t), (1 - p(t)))$$

- If the node is equipped with a split sending a proportion of p_L and p_R to the left and, respectively right, the gained reduction of impurity is:

$$\Delta i(t) = i(t) - p_L i(t_L) - p_R i(t_R)$$

- The best split is the split which maximizes the reduction of impurity
- Other measures of impurity could be used (Gini or Entropy based)

Impurity measures

More generally, for a given node m that defines a region R_M with N_M observations, \hat{p}_{mk} is the observed proportion of cases in class k . The observation at the node will be classified in class $k(m)$ that is the class for which \hat{p}_{mk} is larger. The following impurity measures can be defined:

- Misclassification error:

$$\frac{1}{N_M} \sum_{i \in R_m} I(y_i \neq k(m)) = 1 - \hat{p}_{mk(m)}$$

- Gini index (heterogeneity index):


$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

- Entropy:

$$H = - \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$$

Measures of impurity in two class problems

- for $K = 2$, with p the observed proportion in the second class, these three measures are respectively:
 - $1 - \max(p, 1 - p)$
 - $2p(1 - p) = 2(p - p^2)$
 - $-p \log p - (1 - p) \log(1 - p)$



Avoiding overfitting

- If the overall accuracy is too low we may always make the tree growing further
- The flexibility of the trees would in principle allow for building a perfect classification rule
- A tree that perfectly fits the sample data probably overfits the data: useless for predicting new data, not used for training the tree!
- A useful practice is to evaluate the accuracy of the estimated tree on a test set (out-of-sample).
- Often for Regression and Classification trees the available data are randomly subdivided into three sets:
 - the training set (to grow the tree)
 - the validation set (to prune it)
 - the test set (to evaluate it)
- Evaluation of the quality of the three can be achieved with usual tools for evaluating the prediction (classification) quality: Mean squared prediction errors, confusion matrices, ROC curves (see the R package 'caret')


An example of two class tree

We want to predict now if a life insurance policy is bought using the same covariates

```
TL <- read.csv("TLbin.csv", header=TRUE, sep=",", row.names=1);
attach(TL); set.seed(4321); ind.train <- sample(1:500,300) ;
TL.train <- TL[ind.train,]; TL.test <- TL[-ind.train,]
tree <- rpart(FACEPOS~., data=TL.train); tree
```

```
## n= 300
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 300 132 B (0.5600000 0.4400000)
##      2) INCOME>=23500 240 91 B (0.6208333 0.3791667)
##          4) AGE>=28.5 223 80 B (0.6412556 0.3587444)
##              8) AGE< 68 209 71 B (0.6602871 0.3397129)
##                  16) AGE>=38.5 166 52 B (0.6867470 0.3132530) *
##                      17) AGE< 38.5 43 19 B (0.5581395 0.4418605)
##                          34) AGE< 32.5 15 4 B (0.7333333 0.2666667) *
##                              35) AGE>=32.5 28 13 NB (0.4642857 0.5357143)
##                                  70) MARSTAT=single 7 2 B (0.7142857 0.2857143) *
##                                      71) MARSTAT=not single 21 8 NB (0.3809524 0.6190476)
##                                          142) NUMHH>=3.5 10 4 B (0.6000000 0.4000000) *
##                                              143) NUMHH< 3.5 11 2 NB (0.1818182 0.8181818) *
##          9) AGE>=68 14 5 NB (0.3571429 0.6428571) *
##              5) AGE< 28.5 17 6 NB (0.3529412 0.6470588) *
##      3) INCOME< 23500 60 19 NB (0.3166667 0.6833333)
##          6) EDUCATION>=12.5 27 13 NB (0.4814815 0.5185185)
##              12) MARSTAT=not single 11 2 B (0.8181818 0.1818182) *
##                  13) MARSTAT=single 16 4 NB (0.2500000 0.7500000) *
##                      7) EDUCATION< 12.5 33 6 NB (0.1818182 0.8181818) *
```

The tree



MARS: Multivariate Adaptive Regression Splines


- MARS is an adaptive procedure for regression, and is well suited for highdimensional problems (i.e., a large number of inputs).
- It can be viewed as a generalization of stepwise linear regression or a modification of the CART. This latter approach method to improve the latter's performance in the regression setting.
- a hybrid of MARS called PolyMARS specifically designed to handle classification problems has been also proposed.
- Mars is a semi-parametric method that like CART uses a greedy algorithm and recursively adapt a curve to the regression surface
- At each step it is chosen a couple of basis functions recursively selecting the variable X that is most appropriate and the position of the knot.

- MARS builds models of the form

$$\hat{f}(x) = \sum_{i=1}^k c_i B_i(x)$$

- The model is a weighted sum of basis functions $B_i(x)$. Each c_i is a constant coefficient.
- Each basis function $B_i(x)$ takes one of the following three forms:
 1. a constant
 2. a hinge function. A hinge function has the form $\max(0, x - \text{const})$ or $\max(0, \text{const} - x)$.
MARS automatically selects variables and values of those variables for knots of the hinge functions.
 3. a product of two or more hinge functions. These basis functions can model interaction between two or more variables.

This is an example of a couple of Hinge functions




- Although they might seem quite different, the MARS and CART strategies actually have strong similarities.
- Suppose we take the CART procedure and make the following changes:
 - Replace step functions by the piecewise linear basis functions $I(x - t > 0)$ and $I(x - t \leq 0)$.
 - When a model term is involved in a multiplication by a candidate term, it gets replaced by the interaction, and hence is not available for further interactions.
 - With these changes, the MARS forward procedure is the same as the CART tree-growing algorithm.

An example

```
mod1=earth(V2~V1,data=x,nk=1)
plotmo(mod1,xlab="x",ylab="y", ylim=c(0,1.2)); points(x,pch=20)
```

V2 earth(V2~V1, data=x, nk=1)



An example


```
summary(mod1)

## Call: earth(formula=V2~V1, data=x, nk=1)
##
##               coefficients
## (Intercept)    0.4279076
##
## Selected 1 of 1 terms, and 0 of 1 predictors
## Termination condition: Reached nk 1
## Importance: V1-unused
## Number of terms at each degree of interaction: 1 (intercept only model)
## GCV 0.04290529    RSS 1.202778    GRSq 0    RSq 0
```

An example

```
mod2=earth(V2~V1,data=x,nk=3)  
plotmo(mod2,xlab="x",ylab="y"); points(x,pch=20)
```

V2 earth(V2~V1, data=x, nk=3)



An example

```
summary(mod2)
```

```
## Call: earth(formula=V2~V1, data=x, nk=3)
##
##               coefficients
## (Intercept)      0.7476095
## h(0.176378-V1)  -4.3458394
## h(V1-0.176378)  -0.6146156
##
## Selected 3 of 3 terms, and 1 of 1 predictors
## Termination condition: Reached nk 3
## Importance: V1
## Number of terms at each degree of interaction: 1 2 (additive model)
## GCV 0.00632364    RSS 0.1317425    GRSq 0.852614    RSq 0.8904682
```

Ensemble methods

Combining multiple predictions: Model averaging

- Classification trees can be simple, but often produce noisy (bushy) and weak classifiers
 - **Bagging** (*Breiman, 1996*): Fit many large trees to bootstrap-resampled versions of the training data, and classify by majority vote
 - **Boosting** (*Freund & Shapire, 1996*): Fit many large or small trees to reweighted versions of the training data. Classify by weighted majority vote
 - **Random Forests** (*Breiman 1999*): Fancier version of bagging.
- Note however that the idea of combining multiple predictions or classifications can be used for any technique (i.e., logistic classification, NN, etc.) and it is not limited to trees
- This idea is closely related with model averaging: a strategy for model selection and evaluation of uncertainty in Bayesian analyses

Combining predictions (classifications)

- The idea is to combine the output of different learners for each data point (y, \mathbf{x}) . This help when learners have complementary strengths.
- Suppose training data are available in the form of the p covariates $\mathbf{x} = (x_1, x_2, \dots, x_p)$ and the response is (target) is y . Let $\hat{y}_1 = f_1(\mathbf{x}), \dots, \hat{y}_M = f_M(\mathbf{x})$ be M different predictions (estimates, “experts evaluations”) for the same data point.
- A simple combined vote takes their average

$$f_{comb}(x) = \frac{1}{M} \sum_{m=1}^M f_m(\mathbf{x})$$

In the classification setting for each class k we have a prediction $f_m^k(x, t)$ equal to 0 or 1. Then

$$f_{comb}(x) = \frac{1}{M} \sum_{m=1}^M f_m^k(\mathbf{x})$$

for each class k and $f_{comb}^k(x, t)$ is the proportion of votes for class k . We predict the class with the most number of votes (**majority vote**).

Bagging

- Bagging or bootstrap aggregation averages a given procedure over many samples, to reduce its variance
- A natural way to reduce the variance and increase the prediction accuracy of a statistical learning method is to take many training sets from the population, build a separate prediction model using each training set, and average the resulting predictions.
- Instead, we can
 - bootstrap, by taking repeated samples from the same training data set
 - use the b -th bootstrapped training set to get the prediction $\hat{f}^b(x)$ and
 - average all the predictions, to obtain

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$$

- This is called bagging. In classification problems it uses majority votes.
- Bagging can dramatically reduce the variance of unstable procedure (like trees), leading to improved prediction.
- Bagging averages many trees, produces smoother decision boundaries, reduces the variance, but can slightly increase bias

Out-of-bag

- Using random samples of observations allows the use of the out-of-bag tool, for easy estimation of prediction errors.
- In each bootstrap sample, some of the data of the original training set are excluded.
- On average, each bagged sample makes use of around two-thirds of the observations
- The remaining one-third of the observations not used to fit a given bagged tree are referred to as the out-of-bag (OOB) observations
- For each classifier $\hat{f}^b(\mathbf{x})$ the data of training set that are not in sample can be used as a test set. This will give $B/3$ predictions for the i -th observation.
- Estimate the misclassification error on these data outside the sample used for the fit (out-of-bag), so avoiding cross-classification for large data sets

Random Forest

- A quite popular refinement of bagging. Particularly when bagging trees for which was originally developed. We will describe this version.
- At each tree split, a random sample of m features is drawn, and only those m features are considered for splitting.
- Typically $m = \sqrt{p}$ or $\log_2 p$, where p is the number of features (covariates)
- For each tree grown on a bootstrap sample, the error rate for observations left out of the bootstrap sample is monitored (out-of-bag)
- Random forests tries to improve on bagging by “de-correlating” the trees, and reduce the variance.
- Each tree has the same expectation, so increasing the number of trees does not alter the bias of bagging or random forests.

Variable importance

Random forests can be used to rank the importance of variables in a regression or classification problem.

- For each tree grown in a random forest, calculate number of votes for the correct class in out-of-bag data.
- Now perform random permutation (shuffling) of a predictor's values (let's say variable-k) in the OOB data and then check the number of votes for correct class.
- Subtract the number of votes for the correct class in the variable-k-permuted data from the number of votes for the correct class in the original OOB data.
- The average of this number over all trees in the forest is the raw importance score for variable k. The score is normalized by taking the standard deviation.
- Variables having large values for this score are ranked as more important. It is because if building a current model without original values of a variable gives worse prediction, it means the variable is important.

Boosting

- Designed, initially, exclusively for classification problems
- Idea: Like bagging, but take unequal probability bootstrap samples. Put more weight on observations that are misclassified, to make the classifier work harder on those points. Invention of Freund e Schapire (1997)
- Details
 - Start with equal observation weights $p_i = 1/n$
 - At iteration t , draw a bootstrap sample with the current probabilities p_1, p_2, \dots, p_n , compute the classifier and e_t , the error rate of the classifier on the original sample. Let $\beta_t = e_t / (1 - e_t)$
 - For those points that are classified correctly, decrease their probabilities $p_i = p_i \beta_t$ and normalize them
 - Do this for many (say 1000) iterations.
- At the end, take a weighted vote of the classifications, with weights $\alpha_t = \log(1/\beta_t)$ (more weight on classifiers with lower error).
- Boosting can improve bagging in many instances

Weighting decorrelates the trees, and focuses on regions missed by past trees.